

# Predicting Revenue from Search Engine Advertising Data

MATH2319 - Machine Learning

Course Project

*Ben Cole - s3412349*

*Print Date: 12/06/2019*

## **Contents**

# 1 Phase 1 - Introduction, Cleaning, and Exploration

## 1.1 Outline

The prescribed data set contained advertising metrics provided by a prominent search engine. The data contained several descriptive features pertaining to a range of information. Finally, the target feature was a measure of revenue associated with each of the observations.

The dataset was used to create a supervised machine learning model to predict values for the target feature. Phase 1 of this report contains the introduction, cleaning, and exploration of the dataset. Phase 2 contains the creation, training, and deployment of the machine learning algorithm.

### 1.1.1 Nature of the Data

The below is an excerpt from accompanying documentation about the dataset.

Features in this data set are as follows:

- companyId: Company ID of record (categorical)
- countryId: Country ID of record (categorical)
- deviceType: Device type of record (categorical corresponding to desktop, mobile, tablet)
- day: Day of record (integer between 1 (oldest) and 30 for train, 31 and 35 (most recent) for test)
- dow: Day of week of the record (categorical)
- price1, price2, price3: Price combination for the record set by the company (numeric)
- ad\_area: area of advertisement (numeric)
- ad\_ratio: ratio of advertisement's length to its width (numeric)
- requests, impression, cpc, ctr, viewability: Various metrics related to the record (numeric)
- ratio1, ..., ratio5: Ratio characteristics related to the record (numeric)
- y (target feature): revenue-related metric (numeric)

#### 1.1.1.1 Target Feature

The column/variable **y** was selected as the target feature in the dataset.

#### 1.1.1.2 Descriptive Features

All other columns/variables in the dataset, as outlined above, were chosen as descriptive features.

## 1.2 Data Processing

### 1.2.1 Libraries

The following libraries were used in the below data processing and exploration.

```
library(pacman)                                ## for loading multiple packages

suppressMessages(p_load(character.only = T,
  install = F,
  c("tidyverse", ## thanks Hadley
    "lubridate", ## for handling dates
    "forcats",   ## for categorial variables, not for felines
    "zoo",        ## some data cleaning capabilities
    "lemon",      ## add ons for ggplot
    "rvest",      ## scraping web pages
    "knitr",      ## knitting to RMarkdown
    "kableExtra", ## add ons for knitr tables
    "scales",     ## quick and easy formatting prettynums
    "grid",       ## for stacking ggplots
    "gridExtra",  ## also for stacking ggplots
    "e1071",      ## for skew and kurtosis
    "janitor",    ## cleaning colnames
    "beepR",      ## plays a beep tone
    "mlr",
    "FSelector")))
```

Table 1: Sample of Advertising Data Frame

case_id	companyId	countryId	deviceType	day	dow	price1	price2	price3	ad_area	ad_ratio
165646	43	38	1	24	Monday	0.00	0.00	0.0000	6.5520	0.12363
288	40	77	2	1	Saturday	0.00	0.00	0.0000	0.0001	1.00000
112444	43	56	1	17	Monday	0.59	0.79	1.5610	3.2000	0.31250
2875	43	56	3	1	Saturday	0.00	0.00	0.0000	6.5520	0.12363
95494	43	202	1	14	Friday	0.44	0.80	1.6188	14.4000	0.62500
132826	43	56	3	19	Wednesday	0.00	0.00	0.0000	9.4080	0.83333
211690	43	202	2	30	Sunday	0.00	0.00	0.0000	18.0000	2.00000
73920	95	102	2	12	Wednesday	0.10	0.25	0.2500	0.0001	1.00000
26680	159	144	1	5	Wednesday	0.01	0.21	0.4111	7.5000	0.83333
147331	43	202	2	21	Friday	0.49	1.13	2.2698	6.5520	0.12363
70054	43	13	3	11	Tuesday	0.00	0.00	0.0000	7.5000	0.83333
76481	43	56	2	12	Wednesday	0.00	0.00	0.0000	24.2500	0.25773
111739	43	202	1	17	Monday	0.00	0.00	0.0000	3.2000	0.31250
104359	43	105	2	16	Sunday	0.00	0.00	0.0000	24.2500	0.25773
186218	159	57	3	26	Wednesday	0.12	0.45	0.9040	0.0001	1.00000
60595	43	55	2	10	Monday	0.00	0.00	0.0000	0.0001	1.00000
200946	159	200	2	28	Friday	0.00	0.00	0.0000	0.0001	1.00000
4001	159	197	2	1	Saturday	0.00	0.00	0.0000	0.0001	1.00000
214107	43	234	1	30	Sunday	0.00	0.00	0.0000	0.0001	1.00000
39696	95	38	1	7	Friday	1.18	2.76	5.5200	1.6000	0.15625

### 1.2.2 Loading Data

The prescribed data was made available in comma separated value file format.

```
advertising_train <- read_csv("advertising_train.csv")
```

```
## Parsed with column specification:
```

```
## cols(
##   .default = col_double(),
##   dow = col_character()
## )
```

```
## See spec(...) for full column specifications.
```

```
sample_adv <- sample_n(advertising_train, 20)
```

```
kable_styling(kable(sample_adv[, 1:(ncol(sample_adv)/2)],
  caption = "Sample of Advertising Data Frame",
  font_size = 8.5, latex_options = c("striped"),
  full_width = F)
```

```
kable_styling(kable(sample_adv[, c(1, ((ncol(sample_adv)/2)+1):ncol(sample_adv))],
  caption = "Sample of Advertising Data Frame (cont)",
  font_size = 8.5, latex_options = c("striped"),
  full_width = F)
```

### 1.2.3 Classifying Data

R and dplyr parse data files to guessed data types when loaded. Typically, columns with text are parsed as character type, columns with digits are parsed as numeric, and boolean columns are parsed as logical. Per the above feature definitions, the categorical data was re-classified as factors.

```
advertising_train$companyId <- as.factor(advertising_train$companyId)
```

```
advertising_train$countryId <- as.factor(advertising_train$countryId)
```

Table 2: Sample of Advertising Data Frame (cont)

case_id	requests	impression	cpc	ctr	viewability	ratio1	ratio2	ratio3	ratio4	ratio5	y
165646	58	58	0.2810	0.0517	0.1154	1.0000	0.9483	0.0000	0.6034	0.3793	14.8538462
288	29161	22850	0.0636	0.0003	0.2963	1.0000	0.9737	1.0002	0.0000	0.0000	0.0182470
112444	0	0	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.5344828
2875	0	0	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	1.7000000
95494	1970	1683	0.1141	0.0089	0.7474	0.6423	0.9584	0.0784	0.1040	0.8176	0.8478098
132826	3516	3513	0.1862	0.0043	0.6630	1.0000	0.9004	0.0279	0.2622	0.7111	0.8159578
211690	0	0	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.5029605
73920	0	0	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.1111111
26680	1317	1147	0.0381	0.0026	0.4264	0.3976	0.9991	0.0567	0.1953	0.7489	0.0945137
147331	0	0	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.3098417
70054	68	55	0.1543	0.0182	0.8235	1.0000	0.6909	0.0000	0.5455	0.4545	2.3250000
76481	0	0	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.2605828
111739	794	793	0.1364	0.0038	0.7631	1.0000	0.9420	0.0845	0.1223	0.7932	0.5020378
104359	0	0	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0298861
186218	2189	2175	0.1275	0.0097	0.8476	0.9747	0.8837	0.0441	0.5467	0.4097	1.1674568
60595	0	0	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0500000
200946	719	703	0.0344	0.0028	0.5385	1.0000	0.9260	1.0000	0.0000	0.0000	0.1007018
4001	0	0	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.3640000
214107	280	255	0.1430	0.0157	0.3630	1.0000	0.3882	0.0745	0.4980	0.4275	1.9464883
39696	15663	4181	0.6634	0.0029	0.5794	0.5834	0.8560	0.0440	0.3028	0.6532	0.5599659

```
advertising_train$deviceType <- as.factor(advertising_train$deviceType)
```

```
advertising_train$dow <- as.factor(advertising_train$dow)
```

```
sapply(advertising_train, class)
```

```
##      case_id  companyId  countryId  deviceType      day      dow
## "numeric"   "factor"   "factor"   "factor"   "numeric" "factor"
##      price1    price2    price3    ad_area    ad_ratio  requests
## "numeric"   "numeric" "numeric" "numeric" "numeric" "numeric"
## impression      cpc      ctr viewability      ratio1      ratio2
## "numeric"   "numeric" "numeric" "numeric" "numeric" "numeric"
##      ratio3    ratio4    ratio5      y
## "numeric"   "numeric" "numeric" "numeric"
```

## 1.2.4 Descriptive Statistics

### 1.2.4.1 Numeric Features

The below table outlines basic descriptive statistics about the centre and spread of the data for each of the numeric descriptive features, and numeric target feature. This table indicates that the numeric features each had distributions of different shapes and locations.

```
advertising_train_long_num <- select(advertising_train,
                                   colnames(advertising_train),
                                   -case_id, -countryId,
                                   -companyId, -deviceType,
                                   -dow)
```

```
advertising_train_long_num <- gather(advertising_train_long_num,
                                   key = "Variable",
                                   value = "Value")
```

```
summary_adv_num <- summarise(group_by(advertising_train_long_num,
```

Table 3: Summary Statistics of Numeric Variables

Variable	Mean	Std Dev	Min	Q1	Median	Q3	Max	Number of NA
ad_area	4.724	6.273	0.000	0.000	0.000	7.500	36.000	0.000
ad_ratio	0.923	0.482	0.083	0.833	1.000	1.000	5.000	0.000
cpc	0.178	0.707	0.000	0.000	0.016	0.125	132.534	0.000
ctr	0.033	0.093	0.000	0.000	0.002	0.012	2.000	0.000
day	15.791	8.386	1.000	9.000	16.000	23.000	30.000	0.000
impression	5,585.714	98,713.340	0.000	0.000	99.000	1,058.000	6,100,324.000	0.000
price1	0.438	1.281	0.000	0.000	0.010	0.190	14.690	0.000
price2	0.630	1.482	0.000	0.000	0.090	0.570	63.120	0.000
price3	0.932	1.840	0.000	0.000	0.295	0.986	78.900	0.000
ratio1	0.558	0.447	0.000	0.000	0.750	1.000	1.000	0.000
ratio2	0.491	0.414	0.000	0.000	0.627	0.896	1.027	0.000
ratio3	0.312	0.444	0.000	0.000	0.028	1.000	1.500	0.000
ratio4	0.131	0.240	0.000	0.000	0.000	0.164	1.077	0.000
ratio5	0.188	0.297	0.000	0.000	0.000	0.385	1.200	0.000
requests	8,678.997	122,347.229	0.000	0.000	147.000	1,633.000	6,701,924.000	0.000
viewability	0.378	0.366	0.000	0.000	0.332	0.716	7.000	0.000
y	0.847	1.391	0.000	0.150	0.419	0.959	47.060	0.000

```

      Variable),
      "Mean" = mean(Value, na.rm = T),
      "Std Dev" = sd(Value, na.rm = T),
      "Min" = min(Value, na.rm = T),
      "Q1" = quantile(Value, 0.25, na.rm = T),
      "Median" = median(Value, na.rm = T),
      "Q3" = quantile(Value, 0.75, na.rm = T),
      "Max" = max(Value, na.rm = T),
      "Number of NA" = sum(is.na(Value)))

kable_styling(kable(summary_adv_num,
  digits = 3, format.args = list(nsmall = 3,
                                scientific = F,
                                big.mark = ","),
  caption = "Summary Statistics of Numeric Variables"),
  font_size = 8.5, latex_options = c("striped"),
  full_width = F)

```

#### 1.2.4.2 Categorical and Non-Numeric Features

When examining the frequencies of individual levels of each Categorical (non-numeric) descriptive feature, variability was observed in `companyId`, `countryId`, and `deviceType`. Far less variability in frequencies was observed in `dow`, with Sunday being the only day of the week to return a markedly lower frequency.

```

advertising_train_long_cat <- select(advertising_train,
  countryId,
  companyId, deviceType,
  dow)

advertising_train_long_cat <- gather(advertising_train_long_cat,
  key = "Variable",
  value = "Value")

```

```

## Warning: attributes are not identical across measure variables;
## they will be dropped

```

```

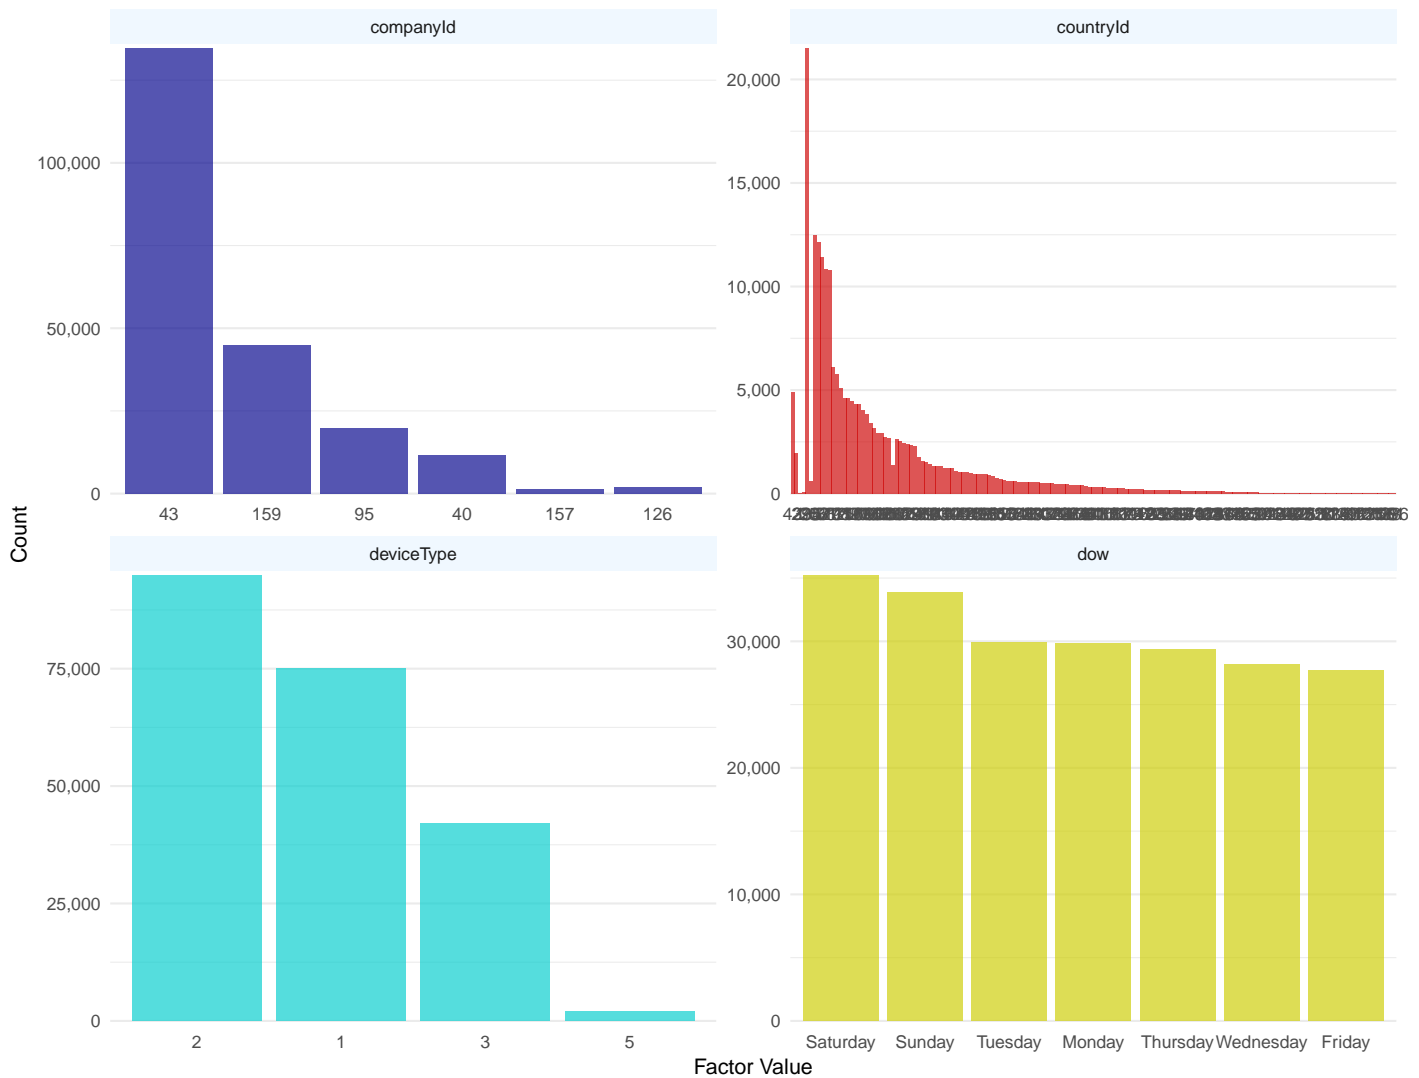
advertising_train_long_cat$Variable <- as.factor(advertising_train_long_cat$Variable)

advertising_train_long_cat$Value <- as.factor(advertising_train_long_cat$Value)

ggplot(advertising_train_long_cat) +
  geom_bar(aes(x = fct_infreq(Value),
               fill = Variable),
           show.legend = F, alpha = 2/3) +
  facet_rep_wrap(~Variable,
                 repeat.tick.labels = T,
                 scales = "free") +
  scale_y_continuous(labels = comma,
                     expand = c(0.01, 0),
                     "Count") +
  scale_x_discrete("Factor Value") +
  scale_fill_manual(values = c("blue4", "red3", "cyan3", "yellow3")) +
  labs(title = "Frequencies of each Value for each Categorical Variable") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
        panel.grid.minor.x = element_blank(),
        strip.background = element_rect(fill = "aliceblue",
                                         colour = NA))

```

## Frequencies of each Value for each Categorical Variable




```
country_labels <- levels(fct_infreq(advertising_train$countryId))[c(seq(1,
                                                                    length(levels(fct_infreq(advertising_train$countryId)))
                                                                    ceiling(length(levels(fct_infreq(advertising_train$countryId))))

ggplot(advertising_train) +
  geom_bar(aes(x = fct_infreq(countryId)),
           fill = "red3", alpha = 2/3) +
  scale_y_continuous(labels = comma,
                    expand = c(0.01, 0),
                    "Count") +
  scale_x_discrete(breaks = country_labels,
                  "countryId") +
  labs(title = "Frequency of observations for each `countryId`",
       subtitle = "(a categorical variable)",
       caption = "labels along x-axis are ID numbers and not numeric/double/ordinal/etc") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
        panel.grid.minor.x = element_blank())
```



Frequency of observations for each `countryId`  
(a categorical variable)



labels along x-axis are ID numbers and not numeric/double/ordinal/etc

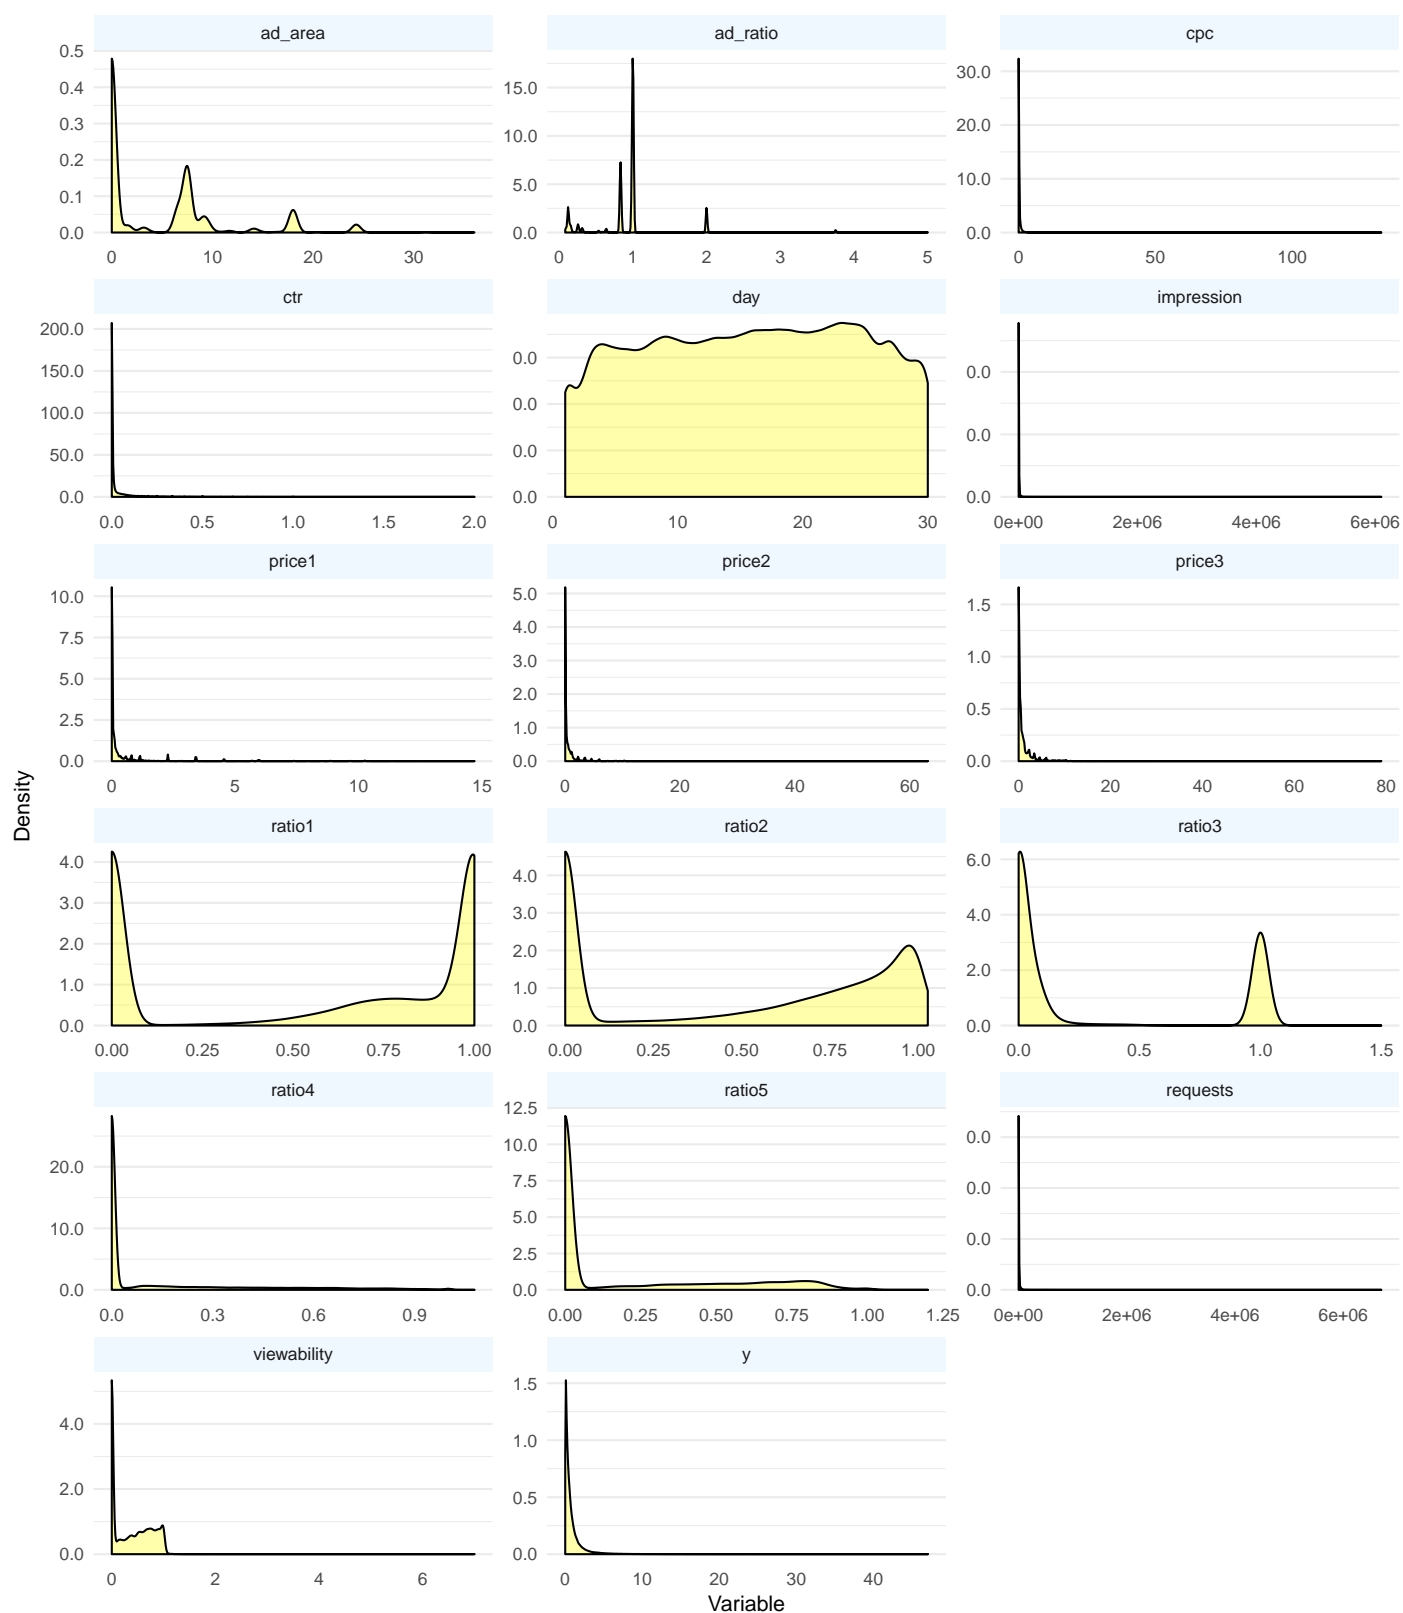
## 1.2.5 Univariate Plots

### 1.2.5.1 Numeric Variables

```
ggplot(advertising_train_long_num) +  
  geom_density(aes(x = Value),  
               fill = "yellow",  
               alpha = 1/3) +  
  facet_rep_wrap(~Variable,  
                 repeat.tick.labels = T,  
                 scales = "free",  
                 ncol = 3) +  
  scale_y_continuous(labels = comma_format(accuracy = 0.1)) +  
  labs(title = "Density Plots of each Numeric Variable",  
       subtitle = "No transformations",  
       x = "Variable",  
       y = "Density") +  
  theme_minimal() +  
  theme(panel.grid.major.x = element_blank(),  
        panel.grid.minor.x = element_blank(),  
        strip.background = element_rect(fill = "aliceblue",  
                                         colour = NA))
```

## Density Plots of each Numeric Variable

No transformations



```
ggplot(advertising_train_long_num) +
  geom_density(aes(x = log(Value)),
    fill = "yellow",
    alpha = 1/3) +
```

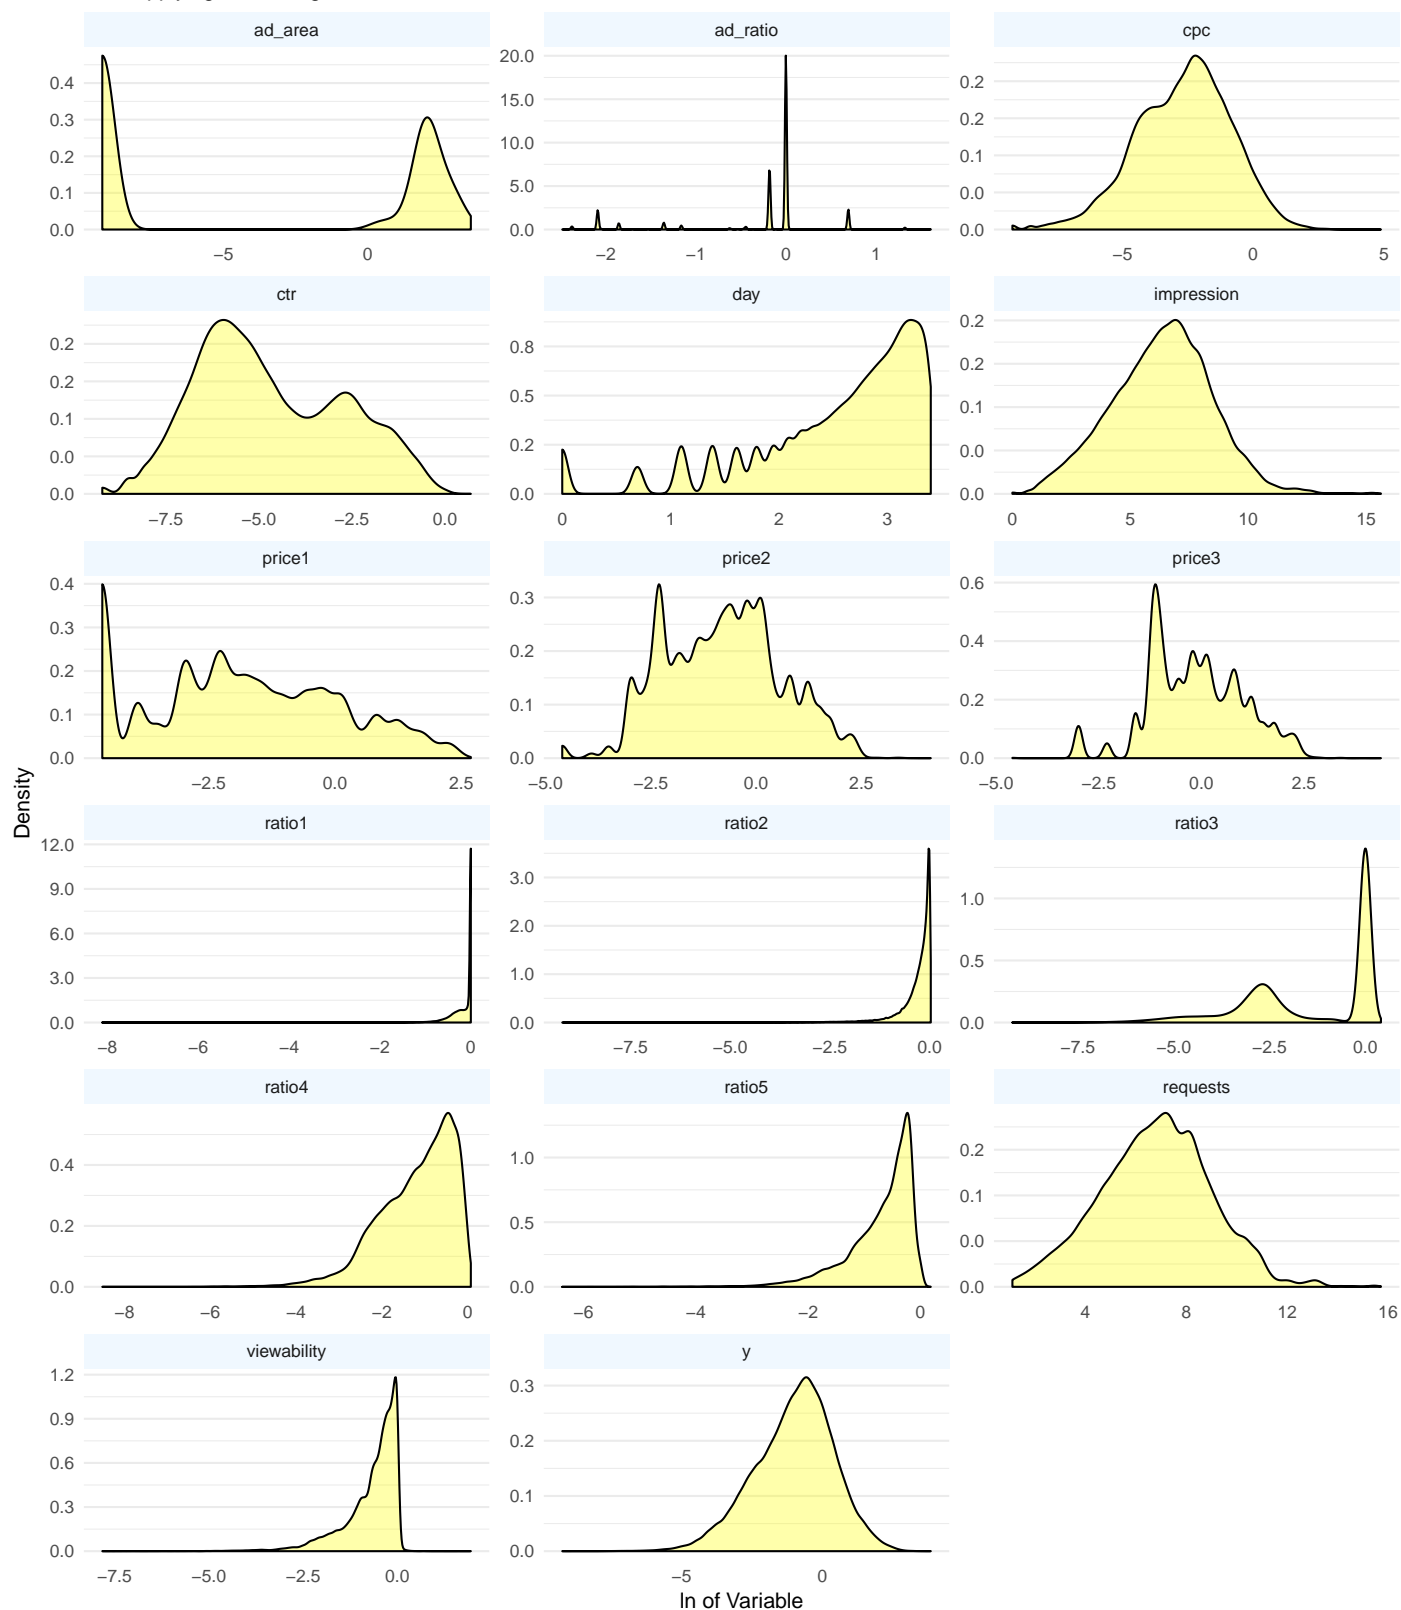
```

facet_rep_wrap(~Variable,
               repeat.tick.labels = T,
               scales = "free",
               ncol = 3) +
scale_y_continuous(labels = comma_format(accuracy = 0.1)) +
labs(title = "Density Plots of each Numeric Variable",
     subtitle = "After applying natural logarithmic transformation",
     x = "ln of Variable",
     y = "Density") +
theme_minimal() +
theme(panel.grid.major.x = element_blank(),
      panel.grid.minor.x = element_blank(),
      strip.background = element_rect(fill = "aliceblue",
                                       colour = NA))

```

```
## Warning: Removed 1213004 rows containing non-finite values (stat_density).
```

Density Plots of each Numeric Variable  
After applying natural logarithmic transformation



### 1.2.5.2 Logarithmic Transformations

It was observed from the plots above that natural logarithmic transformations were applicable for descriptive features `cpc`, `impression`, and potentially `ctr`. Target feature `y` was also suitable for a logarithmic transformation.

Table 4: Sample of advertising\_train Data Frame After Logarithmic Transformations

case_id	companyld	countryld	deviceType	day	dow	price1	price2	price3	ad_area	ad_ratio	requests	impression
146984	43	139	2	21	Friday	0.00	0.00	0.000	0.0001	1.000	0	0
182164	159	98	1	26	Wednesday	0.00	0.00	0.000	7.5000	0.833	5591	5507
1103	43	56	3	1	Saturday	0.00	0.00	0.000	24.2500	0.258	0	0
79692	43	19	2	12	Wednesday	0.00	0.00	0.000	0.0001	1.000	69	64
31797	43	101	1	6	Thursday	0.00	0.00	0.000	6.5520	0.124	382	382
176490	43	103	2	25	Tuesday	0.00	0.00	0.000	0.0001	1.000	47	40
30218	43	127	1	5	Wednesday	0.01	0.14	0.423	7.5000	0.833	0	0
53528	159	144	2	9	Sunday	0.02	0.15	0.283	0.0001	1.000	0	0
123731	159	13	1	18	Tuesday	0.26	0.88	1.762	0.0001	1.000	277	256
52813	43	56	2	9	Sunday	0.00	0.00	0.000	0.0001	1.000	98	98
210661	43	57	3	30	Sunday	0.00	0.00	0.000	7.5000	0.833	12	11
4284	43	59	2	1	Saturday	5.99	5.99	5.988	0.0001	1.000	5572	269
65107	159	57	3	10	Monday	0.14	0.51	1.032	0.0001	1.000	382	381
16126	43	234	1	3	Monday	0.02	0.82	1.642	0.0001	1.000	6180	5115
151997	43	56	2	22	Saturday	0.03	0.92	1.838	24.2500	0.258	268	268
169363	43	202	2	24	Monday	0.82	1.01	2.019	14.0400	0.641	62925	36334
30335	159	144	3	5	Wednesday	0.00	0.00	0.000	0.0001	1.000	0	0
84792	43	179	1	13	Thursday	0.14	0.36	0.723	7.5000	0.833	15909	5552
196327	40	55	1	28	Friday	0.10	0.10	0.200	0.0001	1.000	19061	6880
56897	159	17	1	9	Sunday	0.02	0.07	0.340	0.0001	1.000	13224	7936

```

advertising_train <- mutate(advertising_train,
  "ln_cpc" = log(cpc + 0.005),
  "ln_ctr" = log(ctr + 0.005),
  "ln_impr" = log(impression + 0.005),
  "ln_req" = log(requests + 0.005),
  "ln_y" = log(y + 0.005))

sample_adv <- sample_n(advertising_train, 20)

kable_styling(kable(sample_adv[ , 1 : floor(ncol(sample_adv)/2) ],
  format.args = list(digits = 3),
  caption = "Sample of advertising\\_train Data Frame After Logarithmic Transformations",
  font_size = 8.5, latex_options = c("striped"),
  full_width = F)

kable_styling(kable(sample_adv[ , c(1, seq(from = floor(ncol(sample_adv)/2)+1,
  to = ncol(sample_adv),
  by = 1))],
  format.args = list(digits = 3),
  caption = "Sample of advertising\\_train Data Frame After Logarithmic Transformations",
  font_size = 8.5, latex_options = c("striped"),
  full_width = F)

```

### 1.2.5.3 Comparison of Transformed Features to Normal Curve

As the logarithmic transformation resulted in infinite values, the data frame was trimmed to only include finite values. The finite data frame was then used to calculate the centre and spread of `ln_cpc`, `ln_ctr`, `ln_impr`, `ln_req`, and `ln_y`.

```

finite_cpc <- filter(advertising_train,
  is.finite(ln_cpc))

p_cpc <- ggplot(finite_cpc) +

```

Table 5: Sample of advertising\_train Data Frame After Logarithmic Transformations (cont)

case_id	cpc	ctr	viewability	ratio1	ratio2	ratio3	ratio4	ratio5	y	ln_cpc	ln_ctr	ln_impr	ln_req	ln_y
146984	0.0000	0.0000	0.0000	0.000	0.000	0.0000	0.0000	0.000	0.0250	-5.298	-5.30	-5.30	-5.30	-3.5066
182164	0.0069	0.0045	0.6730	1.000	0.991	0.0750	0.1745	0.751	0.0267	-4.431	-4.66	8.61	8.63	-3.4524
1103	0.0000	0.0000	0.0000	0.000	0.000	0.0000	0.0000	0.000	1.2462	-5.298	-5.30	-5.30	-5.30	0.2241
79692	0.0151	0.0312	0.9412	1.000	1.000	1.0000	0.0000	0.000	0.5760	-3.907	-3.32	4.16	4.23	-0.5430
31797	0.0412	0.0052	0.1671	1.000	0.819	0.4215	0.0366	0.542	0.2507	-3.075	-4.59	5.95	5.95	-1.3639
176490	0.0146	0.1000	1.0000	1.000	0.975	1.0000	0.0000	0.000	0.9388	-3.932	-2.25	3.69	3.85	-0.0578
30218	0.0000	0.0000	0.0000	0.000	0.000	0.0000	0.0000	0.000	0.0143	-5.298	-5.30	-5.30	-5.30	-3.9468
53528	0.0000	0.0000	0.0000	0.000	0.000	0.0000	0.0000	0.000	0.0590	-5.298	-5.30	-5.30	-5.30	-2.7496
123731	0.2843	0.0039	0.7967	0.719	0.648	0.0625	0.4297	0.508	1.3055	-1.240	-4.72	5.55	5.62	0.2704
52813	0.0106	0.1837	0.9032	1.000	0.469	1.0000	0.0000	0.000	1.9472	-4.160	-1.67	4.59	4.59	0.6690
210661	0.0077	0.0909	0.9000	1.000	0.273	0.0000	0.3636	0.545	0.9111	-4.366	-2.34	2.40	2.49	-0.0876
4284	0.0175	0.3123	0.6872	1.000	0.513	1.0037	0.0000	0.000	0.3396	-3.794	-1.15	5.59	8.63	-1.0654
65107	0.1257	0.0184	0.9164	0.992	0.840	0.0210	0.6063	0.373	2.2375	-2.035	-3.76	5.94	5.95	0.8076
16126	0.0440	0.0125	0.5943	0.399	0.877	0.0371	0.6004	0.362	0.5491	-3.016	-4.05	8.54	8.73	-0.5903
151997	0.5060	0.0037	0.7547	0.776	0.888	1.0000	0.0000	0.000	1.8811	-0.671	-4.74	5.59	5.59	0.6345
169363	0.1805	0.0035	0.8172	0.775	0.966	1.0004	0.0000	0.000	0.2521	-1.685	-4.77	10.50	11.05	-1.3582
30335	0.0000	0.0000	0.0000	0.000	0.000	0.0000	0.0000	0.000	0.0267	-5.298	-5.30	-5.30	-5.30	-3.4525
84792	0.0496	0.0045	0.5224	0.704	0.965	0.0657	0.0794	0.855	0.0802	-2.908	-4.66	8.62	9.67	-2.4631
196327	0.0729	0.0015	0.5666	0.989	0.992	0.0811	0.3026	0.616	0.0346	-2.552	-5.04	8.84	9.86	-3.2277
56897	0.0417	0.0009	0.0521	0.429	0.997	0.0931	0.0707	0.836	0.0217	-3.064	-5.13	8.98	9.49	-3.6219

```

geom_density(aes(x = ln_cpc),
              fill = "yellow", alpha = 1/3) +
stat_function(geom = "path", fun = dnorm,
              n = 200, col = "red4", size = 1,
              args = list(mean(finite_cpc$ln_cpc),
                           sd(finite_cpc$ln_cpc))) +
geom_vline(xintercept = mean(finite_cpc$ln_cpc),
           col = "red4", size = 1) +
ylab("Density") +
theme_minimal() +
theme(panel.grid.major.x = element_blank(),
      panel.grid.minor.x = element_blank())

finite_ctr <- filter(advertising_train,
                    is.finite(ln_ctr))

p_ctr <- ggplot(finite_ctr) +
  geom_density(aes(x = ln_ctr),
              fill = "yellow", alpha = 1/3) +
stat_function(geom = "path", fun = dnorm,
              n = 200, col = "red4", size = 1,
              args = list(mean(finite_ctr$ln_ctr),
                           sd(finite_ctr$ln_ctr))) +
geom_vline(xintercept = mean(finite_ctr$ln_ctr),
           col = "red4", size = 1) +
ylab("Density") +
theme_minimal() +
theme(panel.grid.major.x = element_blank(),
      panel.grid.minor.x = element_blank())

finite_impr <- filter(advertising_train,
                     is.finite(ln_impr))

```

```

p_impr <- ggplot(finite_impr) +
  geom_density(aes(x = ln_impr),
    fill = "yellow", alpha = 1/3) +
  stat_function(geom = "path", fun = dnorm,
    n = 200, col = "red4", size = 1,
    args = list(mean(finite_impr$ln_impr),
      sd(finite_impr$ln_impr))) +
  geom_vline(xintercept = mean(finite_cpc$ln_impr),
    col = "red4", size = 1) +
  ylab("Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
    panel.grid.minor.x = element_blank())

finite_req <- filter(advertising_train,
  is.finite(ln_req))

p_req <- ggplot(finite_req) +
  geom_density(aes(x = ln_req),
    fill = "yellow", alpha = 1/3) +
  stat_function(geom = "path", fun = dnorm,
    n = 200, col = "red4", size = 1,
    args = list(mean(finite_req$ln_req),
      sd(finite_req$ln_req))) +
  geom_vline(xintercept = mean(finite_cpc$ln_req),
    col = "red4", size = 1) +
  ylab("Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
    panel.grid.minor.x = element_blank())

finite_y <- filter(advertising_train,
  is.finite(ln_y))

p_y <- ggplot(finite_y) +
  geom_density(aes(x = ln_y),
    fill = "yellow", alpha = 1/3) +
  stat_function(geom = "path", fun = dnorm,
    n = 200, col = "red4", size = 1,
    args = list(mean(finite_y$ln_y),
      sd(finite_y$ln_y))) +
  geom_vline(xintercept = mean(finite_cpc$ln_y),
    col = "red4", size = 1) +
  ylab("Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
    panel.grid.minor.x = element_blank())

ln_vars_title <- textGrob("Logarithmic Transformed Features and Comparison to Normal Curve",
  gp = gpar(fontface = "bold"))

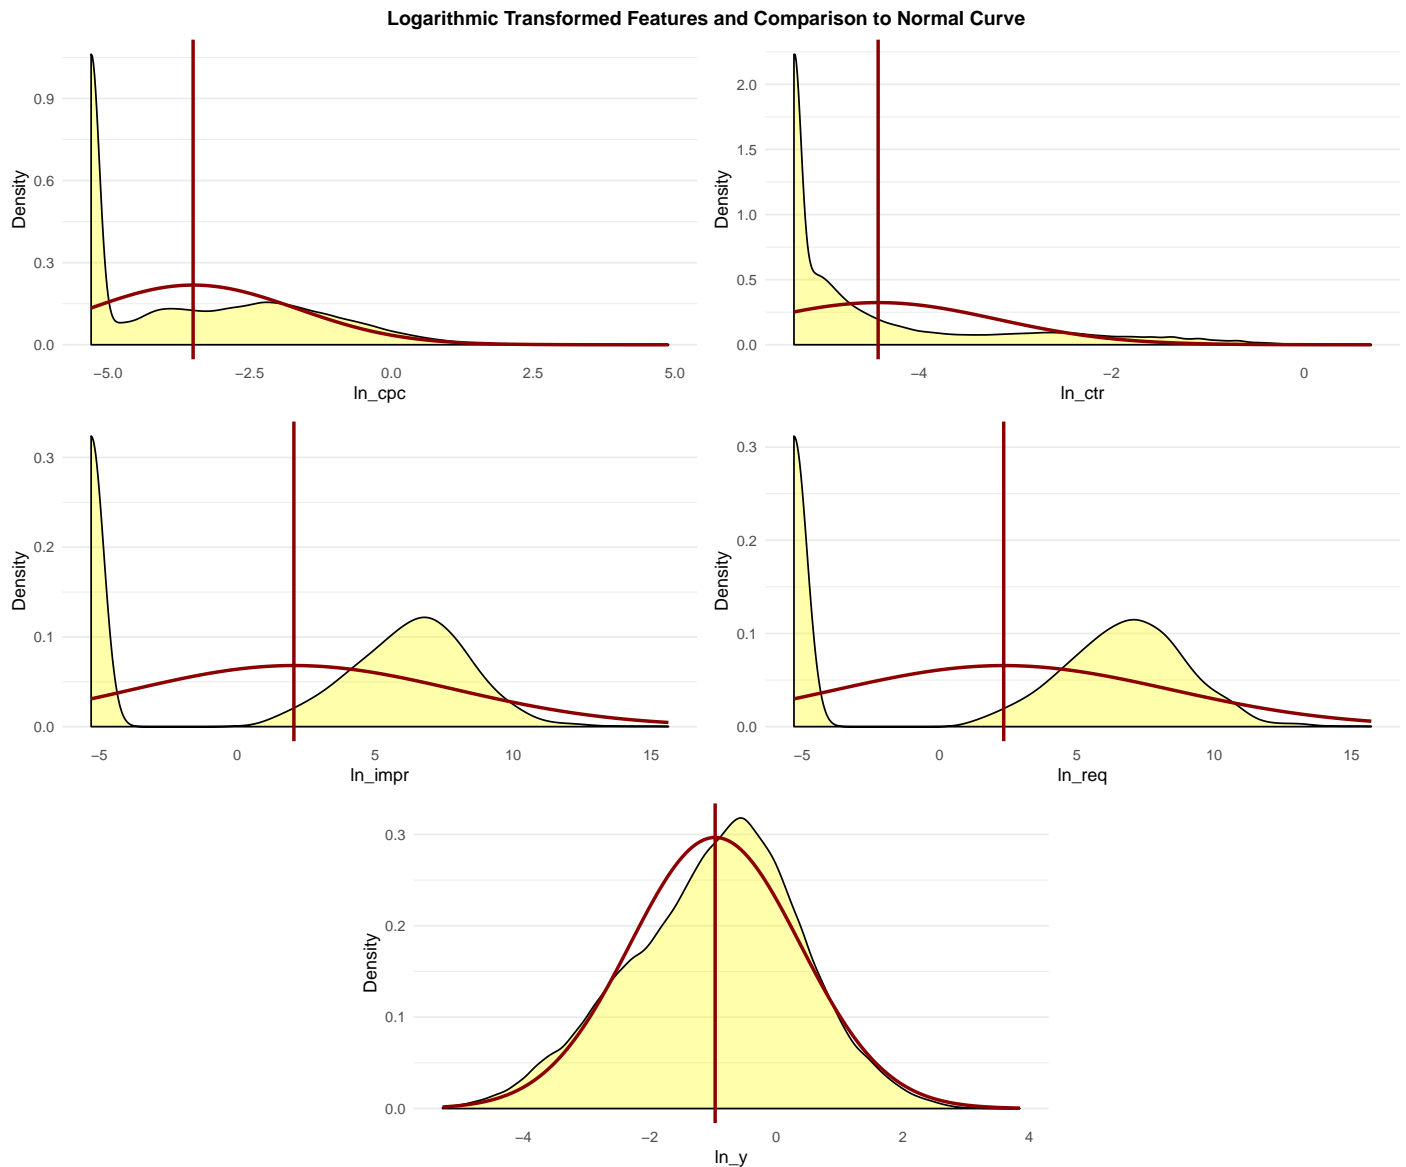
```



```

grid.arrange(top = ln_vars_title,
              p_cpc, p_ctr,
              p_impr, p_req,
              p_y,
              layout_matrix = matrix(c(1,1,2,2,
                                       3,3,4,4,
                                       NA,5,5,NA),
                                    ncol = 4,
                                    byrow = T))

```




The natural logarithmic transformations of `impression` and `requests` clearly approached a normal distribution. The transformed `y` target feature somewhat resembled a normal distribution, albeit less closely as compared to `impression`. Both `cpc` and `ctr` appeared to be bimodal distributions after logarithmic transformation, with `ln_ctr` inarguably so.

### 1.2.6 Multivariate Plots

After transformation, grouping the `ln_ctr` distribution by level within the `companyId` factor revealed several distinct distributions. The distribution for `companyId == 43` still appeared bimodal, which possibly indicated a further dimension of the multivariate relationship.

```
ggplot(advertising_train) +
  geom_density(aes(x = ln_ctr, fill = companyId),
    alpha = 1/3) +
  labs(title = "Density Plots for Logarithmic Transformed `ctr`",
    subtitle = "Grouped by `companyId`",
    y = "Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
    panel.grid.minor.x = element_blank())
```

Density Plots for Logarithmic Transformed `ctr`  
Grouped by `companyId`

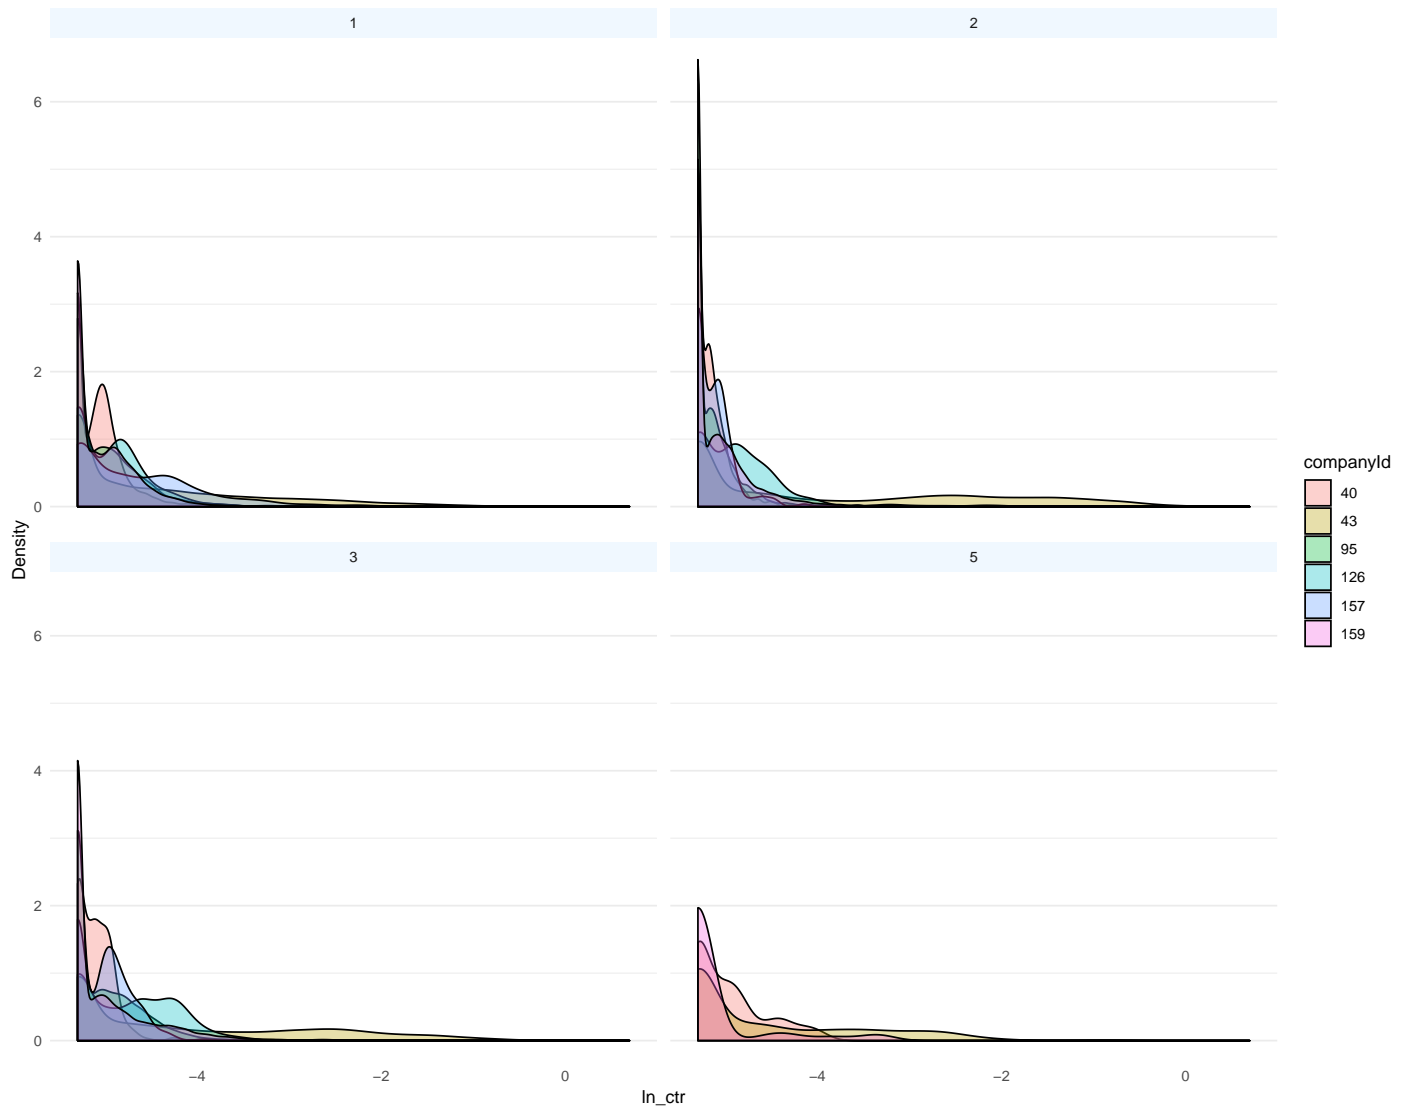


Producing separate density plots for each level within deviceType suggested some trivariate relationship between ln\_ctr, companyId, and deviceType. The effect of facetting by deviceType was particularly apparent when examining companyId == 43, yet it still did not yield Gaussian distributions.

```
ggplot(advertising_train) +
  geom_density(aes(x = ln_ctr, fill = companyId),
    alpha = 1/3) +
  facet_rep_wrap(~deviceType) +
  labs(title = "Density Plots for Logarithmic Transformed `ctr` and each `companyId`",
    subtitle = "Facetted by `deviceType`",
    y = "Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
    panel.grid.minor.x = element_blank(),
```


```
strip.background = element_rect(fill = "aliceblue",
                                colour = NA))
```

Density Plots for Logarithmic Transformed `ctr` and each `companyId`  
Facetted by `deviceType`

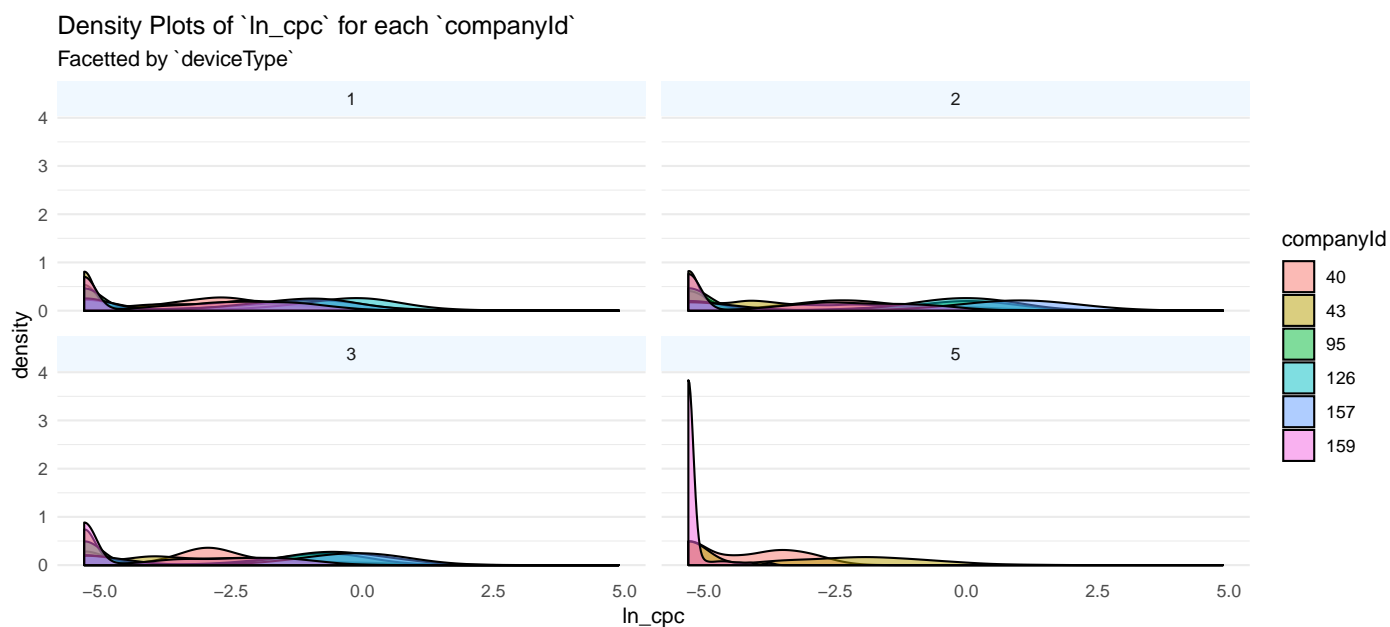


As above for `ln_ctr`, grouping by `companyId` and facetting by `deviceType` revealed a multivariate relationship between aforementioned descriptive features and the transformed `ln_cpc`.

```
ggplot(advertising_train) +
  geom_density(aes(x = ln_cpc, fill = companyId),
              alpha = 1/2) +
  labs(title = "Density Plots of `ln_cpc`",
       subtitle = "Grouped by `companyId`",
       ylab = "Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
        panel.grid.minor.x = element_blank())
```



```
ggplot(advertising_train) +
  geom_density(aes(x = ln_cpc, fill = companyId),
    alpha = 1/2) +
  facet_rep_wrap(~deviceType) +
  labs(title = "Density Plots of `ln_cpc` for each `companyId`",
    subtitle = "Facetted by `deviceType`",
    ylab = "Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
    panel.grid.minor.x = element_blank(),
    strip.background = element_rect(fill = "aliceblue",
    colour = NA))
```



Each of the pricing features, (price1, price2, price3) were not suitably transformed by either logarithmic, square root, or cube root. Logarithmic transformations appeared to spread the data the most, but these transformations considerably diverged from a symmetrical normal distribution. Further grouping by deviceType did not reveal Gaussian distributions.

```

price_trans <- mutate(advertising_train,
                      "ln_price1" = log(price1),
                      "ln_price2" = log(price2),
                      "ln_price3" = log(price3))

p_price1_trans <- ggplot(price_trans) +
  geom_density(aes(x = ln_price1, fill = deviceType),
              alpha = 1/3) +
  labs(y = "Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
        panel.grid.minor.x = element_blank())

p_price2_trans <- ggplot(price_trans) +
  geom_density(aes(x = ln_price2, fill = deviceType),
              alpha = 1/3) +
  labs(y = "Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
        panel.grid.minor.x = element_blank())

p_price3_trans <- ggplot(price_trans) +
  geom_density(aes(x = ln_price3, fill = deviceType),
              alpha = 1/3) +
  labs(y = "Density") +
  theme_minimal() +
  theme(panel.grid.major.x = element_blank(),
        panel.grid.minor.x = element_blank())

price_vars_title <- textGrob("Logarithmic Transformed Price Features",
                             gp = gpar(fontface = "bold"))

grid.arrange(price_vars_title,
              p_price1_trans, p_price2_trans,
              p_price3_trans,
              layout_matrix = matrix(c(1,
                                       2,
                                       2,
                                       2,
                                       3,
                                       3,
                                       3,
                                       4,
                                       4,
                                       4),
                                     ncol = 1,
                                     byrow = T))

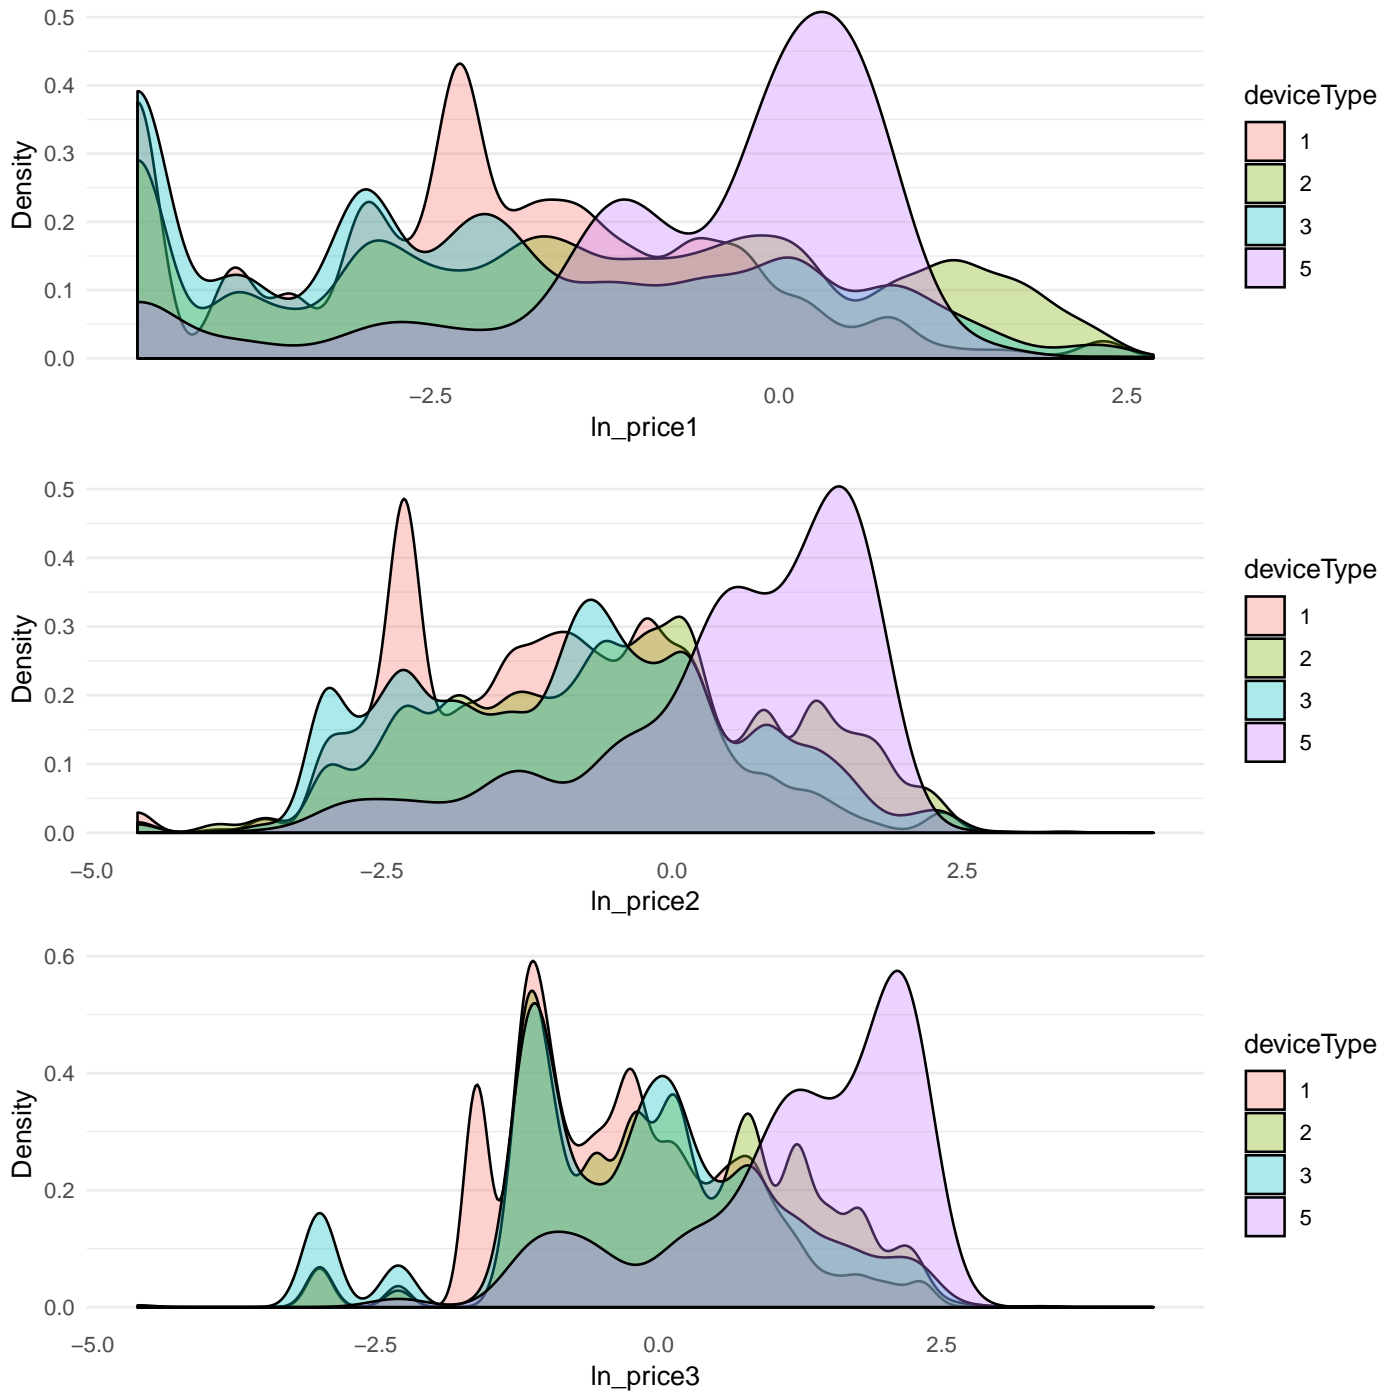
```

```
## Warning: Removed 92892 rows containing non-finite values (stat_density).
```

```
## Warning: Removed 92804 rows containing non-finite values (stat_density).
```

```
## Warning: Removed 92804 rows containing non-finite values (stat_density).
```

### Logarithmic Transformed Price Features



Box-Cox transformations with a range of lambda values also did not convert the price features into distributions that resembled a normal curve.

```
boxcox <- function(x, lambda = 1) {
```

```

(x^(lambda) - 1 /
  (lambda))

}

box_grobs_2 <- list()
box_grobs_higher <- list()

for (i in 1:length(seq(0.025, 0.3, 0.025))) {

  j <- seq(0.025, 0.3, 0.025)[i]

  boxcox_price <- mutate(advertising_train,
                        "bc_price1" = boxcox(x = price1,
                                              lambda = j),
                        "bc_price2" = boxcox(x = price2,
                                              lambda = j),
                        "bc_price3" = boxcox(x = price3,
                                              lambda = j))

  bc_colnames <- colnames(boxcox_price)[str_detect(colnames(boxcox_price), "bc_price")]

  for (k in bc_colnames) {

    m <- which(bc_colnames %in% k)

    box_grobs_2[[m]] <- ggplot(select(boxcox_price,
                                     k, deviceType)) +
      geom_density(aes(x = .data[[k]], fill = deviceType),
                  alpha = 1/3) +
      labs(title = paste("Lambda = ", j)) +
      ylab("Density") + xlab(k) +
      theme_minimal() +
      theme(panel.grid.major.x = element_blank(),
            panel.grid.minor.x = element_blank())

  }

  box_grobs_higher[[i]] <- box_grobs_2

}

density_by_lambda <- list()

for (i in 1:12) {

  density_by_lambda[[i]] <- do.call(what = grid.arrange,
                                    args = list(grobs = box_grobs_higher[[i]],
                                                nrow = 1))

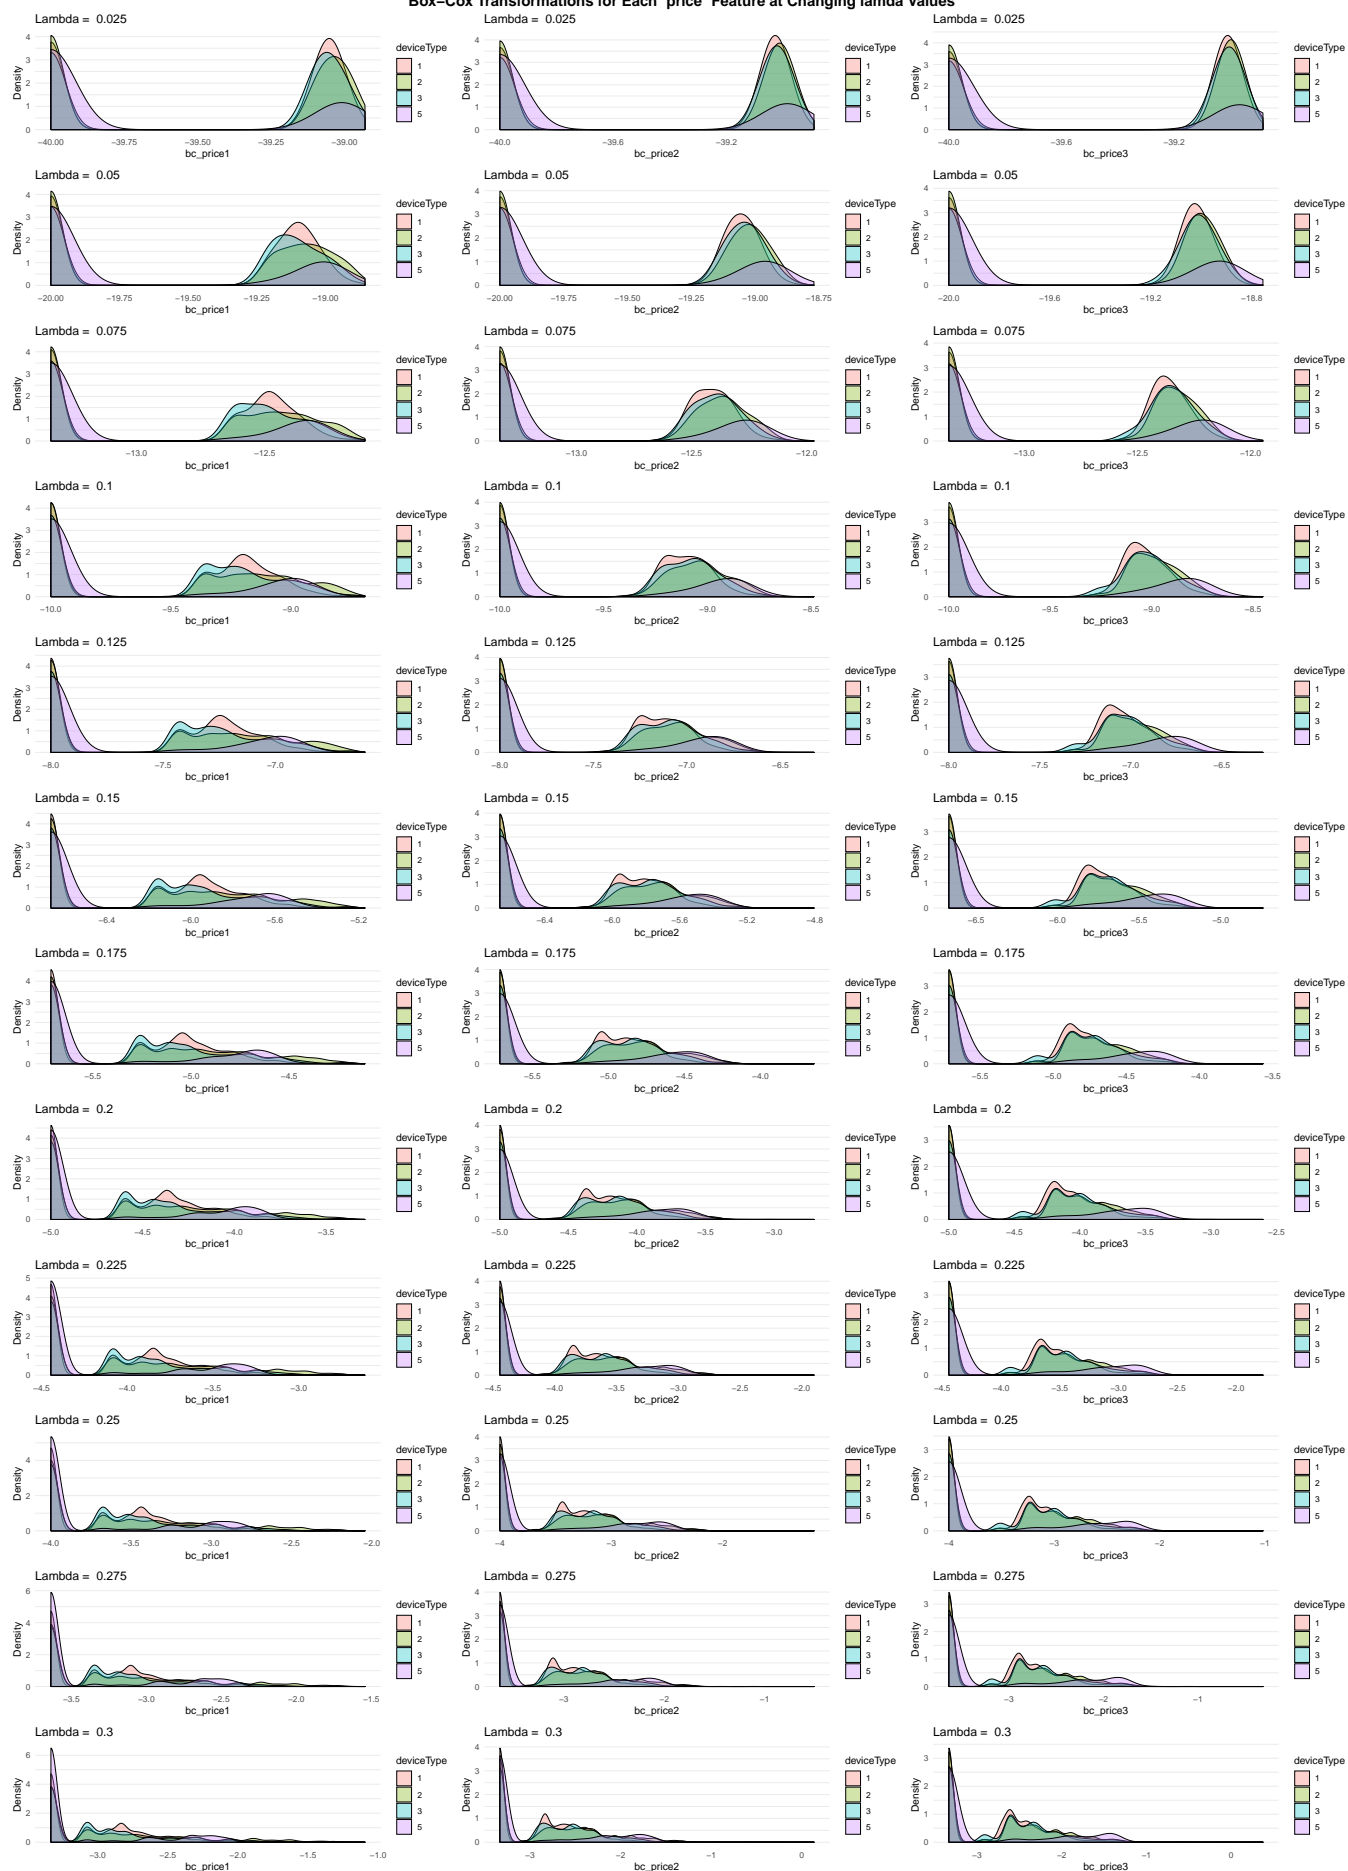
}

```

```
do.call(what = grid.arrange,  
  args = list(grobs = density_by_lambda,  
    top = textGrob("Box-Cox Transformations for Each ``price`` Feature at Changing L  
      gp = gpar(fontsize=16,  
        fontface = "bold")),  
    ncol = 1))
```



# Box-Cox Transformations for Each 'price' Feature at Changing lamda Values



The remaining numeric features (ad\_area, ad\_ratio, day, ratio1, ratio2, ratio3, ratio4, ratio5, and viewability) were not able to be transformed to distributions that approached normal curves via root or logarithmic methods. Despite the accompanying documentation for the prescribed dataset, the ad\_area and day may not strictly be classed as numeric/double variables. Considering the low range, ad\_area could be interpreted as an identifier, and so categorical. The feature day, values 1 - 30, is better interpreted as an ordinal or time value. However, time series forecasting is outside the scope of this project, and so the day feature will be largely ignored from the model and only used for partitioning.

### 1.2.6.1 Data Normalisation

Considering each of the features span differing ranges, both in their raw and transformed applications, it was deemed necessary to normalise each. Normalising the data allowed for more

As outlined in **Fundamentals of Machine Learning**, the below formula was used for normalising the data:

$$a'_i = \left( \frac{a_i - \min(a)}{\max(a) - \min(a)} \right) \times (high - low) + low$$

Where  $a$  is the feature, whether descriptive or target,  $high$  is the highest value in the normalised data range, and  $low$  is the lowest value in the normalised data range. A range of 0 - 1 was chosen, so these values were used for  $low$  and  $high$  respectively.

```
normalise <- function(x) {  
  
  x[is.infinite(x)] <- NA  
  
  (((x - min(x, na.rm = T)) /  
    (max(x, na.rm = T) - min(x, na.rm = T))) * (1 - 0) + 0)  
  
}  
  
num_feats <- select(advertising_train,  
                    case_id,  
                    which(sapply(advertising_train, class)=="numeric"))  
  
for ( i in colnames(num_feats)) {  
  
  newfeat <- paste0("norm_", i)  
  
  advertising_train[[newfeat]] <- normalise(num_feats[[i]])  
  
  advertising_train[[newfeat]][is.na(advertising_train[[newfeat]])] <- advertising_train[[i]]  
  
}  
  
sample_adv <- sample_n(advertising_train, 20)  
  
kable_styling(kable(sample_adv[, 1:floor(ncol(sample_adv)/3)],  
               caption = "Sample of advertising\\_train Data Frame with Normalised Numeric Feat",  
               format.args = list(digits = 2, scientific = F,  
                                   big.mark = ",")),  
              font_size = 8, latex_options = c("striped"),  
              full_width = T)  
  
kable_styling(kable(sample_adv[, c(1,  
                                   seq(from = floor(ncol(sample_adv)/3)*1+1,  
                                       to = floor(ncol(sample_adv)/3)*2,  
                                       by = 1))],  
               caption = "Sample of advertising\\_train Data Frame with Normalised Numeric Feat",  
               format.args = list(digits = 2, scientific = F,  
                                   big.mark = ",")),
```

Table 6: Sample of advertising\_train Data Frame with Normalised Numeric Features (1/3)

case_id	companyid	countryid	deviceType	day	dow	price1	price2	price3	ad_area	ad_ratio	requests	impression	cpc	ctr	viewability
51,925	95	234	3	8	Saturday	0.01	0.74	1.49	7.5000	0.83	1,923	948	1.1486	0.0011	0.21
43,376	43	91	2	7	Friday	0.00	0.00	0.00	0.0001	1.00	24	20	0.0015	0.3000	0.87
30,903	159	200	1	5	Wednesday	0.01	0.27	0.55	0.0001	1.00	243	240	0.0360	0.0042	0.29
34,363	159	110	1	6	Thursday	0.07	0.11	0.34	7.5000	0.83	0	0	0.0000	0.0000	0.00
90,812	43	82	2	14	Friday	5.66	5.66	5.66	0.0001	1.00	200	70	0.0194	0.0286	0.80
186,289	43	12	2	26	Wednesday	1.30	1.72	3.45	24.2500	0.26	0	0	0.0000	0.0000	0.00
110,701	43	36	2	16	Sunday	0.36	1.00	2.01	0.0001	1.00	114	63	0.0342	0.0317	0.98
16,871	159	12	2	3	Monday	0.03	0.15	0.31	0.0001	1.00	0	0	0.0000	0.0000	0.00
115,978	43	12	2	17	Monday	3.42	3.42	3.42	0.0001	1.00	37	32	0.0183	0.2812	0.90
131,004	95	77	2	19	Wednesday	0.01	0.10	0.29	6.5520	0.12	0	0	0.0000	0.0000	0.00
55,896	43	231	2	9	Sunday	4.54	4.54	4.54	0.0001	1.00	327	85	0.0027	0.0941	0.84
195,395	43	13	1	28	Friday	0.00	0.00	0.00	6.5520	0.12	0	0	0.0000	0.0000	0.00
118,799	43	191	1	18	Tuesday	0.80	0.80	0.80	0.0001	1.00	152	42	0.0391	0.0238	1.47
57,434	43	125	2	9	Sunday	0.03	0.07	0.33	6.5520	0.12	0	0	0.0000	0.0000	0.00
89,162	43	113	2	14	Friday	0.00	0.00	0.00	0.0001	1.00	81	33	0.0022	0.1212	0.80
64,224	43	77	3	10	Monday	4.54	4.54	4.54	0.0001	1.00	2,790	798	0.0997	0.0351	0.67
52,335	43	205	1	9	Sunday	0.00	0.00	0.00	0.0001	1.00	0	0	0.0000	0.0000	0.00
130,951	43	95	2	19	Wednesday	0.80	0.80	0.80	0.0001	1.00	480	108	0.1444	0.0185	0.83
158,670	43	179	2	23	Sunday	5.96	5.96	5.96	0.0001	1.00	63,256	1,881	0.0288	0.4588	0.84
175,709	43	179	2	25	Tuesday	0.12	0.31	0.61	6.5520	0.12	0	0	0.0000	0.0000	0.00

Table 7: Sample of advertising\_train Data Frame with Normalised Numeric Features (2/3)

case_id	ratio1	ratio2	ratio3	ratio4	ratio5	y	ln_cpc	ln_ctr	ln_impr	ln_req	ln_y	norm_caseid	norm_day	norm_price1	norm_price2	norm_price3
51,925	0.86	0.78	0.0084	0.55	0.44	0.452	0.14	-5.10	6.9	7.6	-0.78	0.242	0.241	0.00068	0.0117	0.0189
43,376	1.00	0.95	1.0000	0.00	0.00	0.052	-5.04	-1.19	3.0	3.2	-2.86	0.203	0.207	0.00000	0.0000	0.0000
30,903	0.71	0.99	0.0583	0.10	0.84	0.167	-3.19	-4.69	5.5	5.5	-1.76	0.144	0.138	0.00068	0.0043	0.0069
34,363	0.00	0.00	0.0000	0.00	0.00	0.101	-5.30	-5.30	-5.3	-5.3	-2.24	0.160	0.172	0.00477	0.0017	0.0044
90,812	1.00	0.81	1.0000	0.00	0.00	0.033	-3.71	-3.39	4.2	5.3	-3.28	0.424	0.448	0.38530	0.0897	0.0717
186,289	0.00	0.00	0.0000	0.00	0.00	0.238	-5.30	-5.30	-5.3	-5.3	-1.41	0.870	0.862	0.08850	0.0272	0.0437
110,701	0.65	1.00	1.0000	0.00	0.00	0.440	-3.24	-3.30	4.1	4.7	-0.81	0.517	0.517	0.02451	0.0158	0.0254
16,871	0.00	0.00	0.0000	0.00	0.00	0.338	-5.30	-5.30	-5.3	-5.3	-1.07	0.079	0.069	0.00204	0.0024	0.0039
115,978	0.97	0.88	1.0000	0.00	0.00	5.310	-3.76	-1.25	3.5	3.6	1.67	0.542	0.552	0.23281	0.0542	0.0433
131,004	0.00	0.00	0.0000	0.00	0.00	0.175	-5.30	-5.30	-5.3	-5.3	-1.72	0.612	0.621	0.00068	0.0016	0.0037
55,896	1.00	0.42	1.0000	0.00	0.00	0.372	-4.87	-2.31	4.4	5.8	-0.98	0.261	0.276	0.30905	0.0719	0.0575
195,395	0.00	0.00	0.0000	0.00	0.00	1.136	-5.30	-5.30	-5.3	-5.3	0.13	0.913	0.931	0.00000	0.0000	0.0000
118,799	1.00	0.26	0.0476	0.00	0.95	0.398	-3.12	-3.55	3.7	5.0	-0.91	0.555	0.586	0.05446	0.0127	0.0101
57,434	0.00	0.00	0.0000	0.00	0.00	0.033	-5.30	-5.30	-5.3	-5.3	-3.28	0.268	0.276	0.00204	0.0011	0.0042
89,162	1.00	0.48	1.0000	0.00	0.00	0.038	-4.93	-2.07	3.5	4.4	-3.15	0.416	0.448	0.00000	0.0000	0.0000
64,224	0.98	0.64	0.0113	0.69	0.30	1.563	-2.26	-3.22	6.7	7.9	0.45	0.300	0.310	0.30905	0.0719	0.0575
52,335	0.00	0.00	0.0000	0.00	0.00	0.259	-5.30	-5.30	-5.3	-5.3	-1.33	0.244	0.276	0.00000	0.0000	0.0000
130,951	0.99	0.27	1.0000	0.00	0.00	0.672	-1.90	-3.75	4.7	6.2	-0.39	0.612	0.621	0.05446	0.0127	0.0101
158,670	0.77	0.81	1.0005	0.00	0.00	0.444	-3.39	-0.77	7.5	11.1	-0.80	0.741	0.759	0.40572	0.0944	0.0755
175,709	0.00	0.00	0.0000	0.00	0.00	0.031	-5.30	-5.30	-5.3	-5.3	-3.33	0.821	0.828	0.00817	0.0049	0.0078

```
font_size = 8, latex_options = c("striped"),
full_width = T)
```

```
kable_styling(kable(sample_adv[, c(1,
                                seq(from = floor(ncol(sample_adv)/3)*2+1,
                                      to = floor(ncol(sample_adv)/3)*3,
                                      by = 1))],
              caption = "Sample of advertising\\_train Data Frame with Normalised Numeric Features",
              format.args = list(digits = 2, scientific = F,
                                  big.mark = ",")),
font_size = 8, latex_options = c("striped"),
full_width = T)
```